

## 1. GRPO Demo - Gemma3 1B (Baseline)

File: grpo-demo-gemma3-1b.ipynb

### Model Setup & Architecture

- **Model:** Gemma 3 1B-IT
- **Method:** Group Relative Policy Optimization (GRPO)
- **LoRA Configuration:** Rank=64, Alpha=64.0
- **Framework:** JAX/Tunix/Flax, TPU v5e-8
- **Sharding:** FSDP + Tensor Parallelism across 4 devices `[(1, 4), ("fsdp", "tp")]`

### Training Pipeline

- **Dataset:** GSM8K (grade school math) via HuggingFace
- **Data Format:** Structured XML-style tags `<reasoning>`, `</reasoning>`, `<answer>`, `</answer>`
- **Preprocessing:** Extracts answers from `####` delimiter format
- **Batch Size:** 4 (micro-batch)
- **Total Batches:** 3,738
- **Epochs:** 1

### Reward Function & RL Components

**Multi-Component Reward System** (4 reward functions):

1. `match format exactly`: +3.0 points for perfect XML structure
2. `match format approximately`: ±0.5 per tag (penalizes duplicates)
3. `check answer`:
  - +3.0 for exact match
  - +1.5 for whitespace-normalized match
  - +0.5/+0.25 for answers within 10%/20% ratio
  - -1.0 penalty for wrong answers
4. `check numbers`: +1.5 for extracted numerical correctness

### GRPO Hyperparameters:

- Generations per prompt (G): 4
- KL penalty ( $\beta$ ): 0.08
- PPO-style clipping ( $\epsilon$ ): 0.2
- Iterations per batch ( $\mu$ ): 1

### Configurations

- **Optimizer:** AdamW (LR=3e-6,  $\beta_1$ =0.9,  $\beta_2$ =0.99)
- **LR Schedule:** 10% warmup → Cosine decay to 0
- **Gradient Clipping:** 0.1 (critical for KL stability)

- **Weight Decay:** 0.1
- **Generation:** Temperature=0.9, top\_k=50, top\_p=1.0 (high diversity during training)
- **Checkpoint:** Every 500 steps, keep 4 latest

### What's Unique

- **Comprehensive reward shaping** with tolerance for partial correctness
- **Format enforcement** through dual exact/approximate matching
- **Ratio-based scoring** (0.9-1.1 range) for near-correct numerical answers
- Standard GRPO implementation following DeepMind's algorithm

### Strengths

- Well-documented baseline
- Robust multi-reward system with graceful degradation
- Reproducible HuggingFace dataset pipeline
- Memory-efficient GRPO (no separate value model)

### Weaknesses

- No trajectory-level rewards (only token-level)
  - High temperature (0.9) may hurt convergence
  - Single epoch limits learning potential
  - No advanced techniques (e.g., curriculum learning, self-consistency)
- 

## 2. "Efficiency Build" - Gemma3 1B (Memory-Optimized)

**Files:** start-with-gemma3-1b-it-tutorial.ipynb, [start-with-gemma3-1b-it-tutorial \(1\).ipynb](#)

### Model Setup & Architecture

- **Model:** Gemma 3 1B-IT
- **Key Difference:** Resource-constrained adaptation for Kaggle TPU v3-8
- **LoRA Configuration:** Rank=32, Alpha=32.0 (halved from baseline!)
- **Batch Size:** 2 (reduced from 4)

### Training Pipeline

- Identical dataset/preprocessing to baseline
- **Custom auto-login logic** for Kaggle secrets (UX improvement)
- Same XML tag structure

### Reward Function & RL Components

- **Identical** 4-reward system to baseline
- Same GRPO hyperparameters ( $\beta=0.08$ ,  $\varepsilon=0.2$ ,  $G=4$ )

## Configurations

### Memory Optimizations:

- LoRA rank: 64 → **32** (-50% trainable params)
- Batch size: 4 → **2** (-50% memory)
- Maintains same LR/optimizer settings

### What's Unique

- **Explicitly designed for OOM prevention** on lower-tier TPUs
- **Documented trade-off philosophy:** "Efficiency over capacity"
- Auto-login helper for Kaggle environment
- Claims "significant performance gains even with limited compute budgets"

### Strengths

- Practical for resource-constrained users**
- Preserves training stability with reduced capacity
- Better UX with auto-login logic
- Demonstrates LoRA compression effectiveness

### Weaknesses

- Lower model capacity** (rank 32 vs 64) may hurt complex reasoning
  - Smaller batches increase gradient variance
  - No unique algorithmic contributions
  - Unverified claims about "performance gains" (no metrics provided)
- 

## 3. Supervised Fine-Tuning (Full) - Gemma3 1B

File: supervised-fine-tuning-full.ipynb

### Model Setup & Architecture

- **Model:** Gemma 3 1B
- **Method: Full fine-tuning** (NOT LoRA! All 1B parameters trainable)
- **Framework:** JAX/Tunix/Flax, TPU v5e-8
- **Sharding:** `(8, 1)` - FSDP across all 8 TPU cores

### Training Pipeline

#### Dataset Cleaning (Critical differentiator):

```
def clean_gsm8k_content(text):
    # Removes GSM8K calculation artifacts like <>10+5=15>
```

```
return text.replace("<<", "(").replace(">>", ")")
```

- **Applies cleaning BEFORE training** to reduce noise
- Uses "STRICT System Prompt" emphasizing format compliance

## Data Processing:

- Max sequence length: 2048
- Loss mask: Only active for model response (not prompt)
- Batch size: 2 (micro),  $2 \times 8 \times 4 = 64$  (effective global batch with grad accumulation)

## Reward Function & RL Components

N/A - This is supervised learning, not RL!

- Standard cross-entropy loss
- No reward functions
- Loss computed only on completion tokens (prompt is masked)

## Configurations

- **Optimizer**: AdamW (LR=2e-5,  $\beta_1=0.9$ ,  $\beta_2=0.999$ ,  $\epsilon=1e-8$ )
- **LR Schedule**: 50-step warmup → Cosine decay to  $0.1 \times LR$
- **Gradient Accumulation**: 4 steps (effective batch=64)
- **Gradient Clipping**: 1.0 (10× looser than GRPO)
- **Weight Decay**: 0.01 (10× less aggressive)
- **Epochs**: 10 (far more than GRPO's 1)
- **Max Steps**: 1,170 (117 batches × 10 epochs)

## Post-Training Inference:

- **Majority Voting** with k=1 sample (self-consistency)
- Temperature=0.7 for diverse reasoning paths
- Robust answer extraction (XML → LaTeX boxed → text patterns → last number)

## What's Unique

- 🔥 **Only full fine-tuning approach** (vs parameter-efficient LoRA)
- 🔥 **Data cleaning pipeline** to remove GSM8K artifacts
- 🔥 **Majority voting evaluation** (self-consistency)
- 🔥 **Multi-pattern answer extraction** (most robust)
- 🔥 **10 epochs** vs 1 for RL approaches

## Strengths

- ✓ **Full model capacity** utilization
- ✓ **Data quality focus** (cleaning artifacts)
- ✓ **Proper evaluation methodology** (voting, robust extraction)
- ✓ **Reproducible** with clear TPU verification steps
- ✓ **Best for datasets with clean supervision** (GSM8K has gold reasoning traces)

## Weaknesses

- ✗ **No parameter efficiency** (trains all 1B params)
- ✗ **Requires high-quality reasoning traces** in training data
- ✗ **Doesn't learn reward optimization** (just imitation)
- ✗ **10x slower** than single-epoch GRPO
- ✗ **Risk of overfitting** with 10 epochs on same 7.5k examples

## 4. TT-Tunix-on-TPU (Incomplete Attempt)

File: tt-tunix-on-tpu.ipynb

### Model Setup & Architecture

- **Model:** Gemma2 2B-IT
- **Attempted Method:** GRPO (based on Windmaple's code)
- **LoRA Configuration:** Rank=64, Alpha=64.0

### Training Pipeline

- **Data Source:** Kaggle CSV (The Devastator's GSM8K)
- Setup code for Gemma2 2B checkpoint loading

### Reward Function & RL Components

- Copied baseline 4-reward system
- No modifications

### Configurations

- Batch size: 2
- Standard GRPO hyperparameters

### What's Unique

✗ **Incomplete submission** - Training failed at LoRA policy creation:

TypeError: set\_metadata takes either 1 argument or 1 or more keyword arguments,

got args=('sharding\_names', ('fsdp', None))

- Demonstrates **common pitfall**: Flax/JAX version incompatibility with Tunix

### Strengths

- ✓ Attempted larger model (2B)
- ✓ Documented failure (useful for debugging)

### Weaknesses

- ✗ Non-functional code
- ✗ No unique contributions
- ✗ Version mismatch issues
- ✗ Abandoned midway

## 5. Tunix-Hack-Winner: Trajectory Reward Training

File: tunix-hack-winner-trajectory-reward-training.ipynb

### Model Setup & Architecture

- Model: Gemma2 2B-IT
- Method: GRPO with **trajectory-level rewards**
- LoRA Configuration: Rank=64, Alpha=64.0
- Framework: JAX/Tunix, TPU v4/v5

### Training Pipeline

- Dataset: Kaggle CSV (GSM8K main\_train.csv)
- Batch Size: 2 (safe for V4 TPUs)
- Batches: 5,000 (vs 3,738 in baseline)
- Standard preprocessing

### Reward Function & RL Components

🔥 COMPETITIVE ADVANTAGE: Trajectory Reward

```
def trajectory_reward(prompts, completions, answer, **kwargs):  
  
    # Reasoning quality (50% weight initially, tuned to 60%)  
  
    reasoning = re.search(r'<reasoning>(.*)?</reasoning>', comp, re.DOTALL)  
  
    r_score = 0.5 if reasoning and len(reasoning.group(1).split()) > 20 else 0.0  
  
  
    # Answer correctness (50% weight initially, tuned to 40%)  
  
    ans = re.search(r'<answer>(.*)?</answer>', comp, re.DOTALL)  
  
    a_score = 1.0 if ans and ans.group(1).strip() == ref else 0.0  
  
  
    # Weighted combination: 60% reasoning + 40% answer
```

`reward = (0.6 * r_score + 0.4 * a_score) * 3.0`

### Key Innovation:

- **Holistic evaluation:** Entire completion scored as a unit
- **Reasoning length threshold:** 20+ words required
- **Weighted objectives:** 60% process, 40% outcome
- **Scaled rewards:** 3x multiplier for final score

### Configurations

- **Optimizer:** AdamW (LR=5e-5, higher than baseline's 3e-6)
- **GRPO:** G=2 generations (vs 4 in baseline),  $\beta=0.08$ ,  $\epsilon=0.2$
- **Rollout:** 768 max tokens (vs 512)
- **Batch Processing:** V4-safe settings

### What's Unique

- 🏆 **Claimed winner approach**
- 🔥 **Trajectory-level reward** vs token-level
- 🔥 **Explicit reasoning quality metric** (word count threshold)
- 🔥 **60/40 weighting** between process and outcome
- 🔥 **Higher learning rate** (5e-5 vs 3e-6)
- 🔥 **Fewer generations** (2 vs 4) for efficiency

### Strengths

- ✓ **Process-oriented reward design** (aligns with hackathon goals!)
- ✓ **Scalable to larger models** (2B vs 1B)
- ✓ **Efficient generation** (G=2 reduces compute)
- ✓ **Transparent reasoning evaluation** (simple word count)
- ✓ **Reward scaling** prevents vanishing signals

### Weaknesses

- ✗ **Crude reasoning metric** (word count ≠ quality)
- ✗ **No citation counting** or logical step verification
- ✗ **Fixed threshold (20 words)** not adaptive
- ✗ **Higher LR** may cause instability
- ✗ **No reproducibility verification** (no training logs shown)

**Cross-Project Comparison Table**

Aspect	GRPO Baseline	Efficiency Build	SFT Full	TT-Tunix	Trajectory Winner
<b>Model</b>	Gemma3 1B	Gemma3 1B	Gemma3 1B	Gemma2 2B	Gemma2 2B
<b>Method</b>	GRPO	GRPO	Supervised	GRPO (failed)	GRPO
<b>LoRA Rank</b>	64	<b>32</b>	N/A (full)	64	64
<b>Batch Size</b>	4	<b>2</b>	2 ( $\times 8 \times 4 = 64$ eff.)	2	2
<b>Epochs</b>	1	1	<b>10</b>	-	1
<b>Reward Type</b>	Multi-component	Multi-component	N/A	-	<b>Trajectory</b>
<b>Data Cleaning</b>	✗	✗	✓	✗	✗
<b>Reasoning Metric</b>	Format only	Format only	N/A	-	<b>Word count</b>
<b>Evaluation</b>	Greedy	Greedy	<b>Majority vote</b>	-	Unknown
<b>Reproducible</b>	✓	✓	✓	✗	⚠

## Alignment with Hackathon Goals

**Goal:** "Teach model to show its work" (reasoning transparency + accuracy + reproducibility)

### Rankings:

1. 🥇 **Trajectory Winner** - Explicitly rewards reasoning quality (60% weight), transparency-first
2. 🥈 **SFT Full** - Comprehensive pipeline with evaluation rigor, but relies on supervision
3. 🥉 **GRPO Baseline** - Solid foundation, format-enforcing, fully reproducible
4. **Efficiency Build** - Practical but no innovation
5. **TT-Tunix** - Non-functional

## Strengths & Weaknesses Summary

### Strongest Ideas:

1. **Trajectory-level rewards** (Winner) - Holistic reasoning evaluation
2. **Data cleaning** (SFT) - Removes artifacts for cleaner learning
3. **Majority voting** (SFT) - Self-consistency improves reliability
4. **Multi-reward shaping** (Baseline) - Graceful handling of partial correctness
5. **LoRA compression** (Efficiency) - Practical accessibility

### Methods Likely to Fail:

1. **Word count as reasoning quality** - Easily gamed, no semantic understanding

- 2. **Single-epoch RL** - May underfit complex reasoning patterns
- 3. **High temperature (0.9)** - Reduces convergence stability
- 4. **No trajectory rewards in baseline** - Misses process-oriented learning

#### Most Reproducible:

- 1. **SFT Full** - Complete pipeline with TPU verification steps
- 2. **GRPO Baseline** - Standard Tunix tutorial format
- 3. **Efficiency Build** - Documented resource constraints

#### Easiest to Reproduce:

**GRPO Baseline** - Uses default Tunix APIs, HuggingFace datasets, no custom logic

## Recommended Hybrid Approach

Combining the best ideas:

```
# Model: Gemma3 1B (SFT-pretrained)

# Phase 1: SFT with cleaned data (3 epochs)

# Phase 2: GRPO with hybrid rewards
```

```
def hybrid_reward(prompts, completions, answer, **kwargs):
    rewards = []
    for comp, ref in zip(completions, answer):
        # 1. Format enforcement (Baseline idea)
        format_score = 3.0 if match_format.search(comp) else 0.0

        # 2. Reasoning depth (Winner idea, improved)
        reasoning = re.search(r'<reasoning>(.*)?</reasoning>', comp, re.DOTALL)
        if reasoning:
            steps = len(re.findall(r'\n\d+\.\|\n-', reasoning.group(1))) # Count numbered steps
            words = len(reasoning.group(1).split())
```

```
depth_score = min(steps * 0.3 + (words > 50) * 0.5, 2.0) # Cap at 2.0
else:
    depth_score = 0.0

# 3. Answer accuracy (Baseline idea)
ans = re.search(r'<answer>(.*)?</answer>', comp, re.DOTALL)
if ans:
    try:
        acc_score = 3.0 if float(ans.group(1)) == float(ref) else 0.0
    except:
        acc_score = 0.0
    else:
        acc_score = 0.0

# Weighted: 30% format + 40% reasoning + 30% answer
total = 0.3*format_score + 0.4*depth_score + 0.3*acc_score
rewards.append(total)

return rewards

# Config: LoRA rank=64, batch=4, LR=3e-6, epochs=2
# Evaluation: Majority voting (k=5)
```

### Why this works:

- SFT initializes reasoning structure
- GRPO refines quality with process-oriented rewards
- Step counting > word counting for depth
- Balanced weights across format/process/outcome
- Self-consistency for robust evaluation

---

**Final Recommendation:** The **Trajectory Winner** has the right philosophy (process over outcome), but implementation needs refinement. The **SFT Full** approach provides the most complete pipeline for reproducibility. A hybrid SFT→GRPO approach combining both would likely dominate.